

Programming Assignment Multi-Threading 1 and Debugging 1

Due Date: **Friday, May 13 @ 11:59pm**

Overview:

The purpose of this mini-assignment is to briefly introduce you to parallel programming and multi-threading; covering aspects of implementation, benefits and the tradeoffs involved. Provided with skeleton code, you will fill in the functionality where necessary to initialize an array with decoded values, calculate the sum of the squares, min, and max of the values in the array, and calculate the square root of the values, min, and max of the values in the array -- first with sequential algorithms and then using OpenMP directives to parallelize portions of the algorithms. The program will also record and display the execution time of the initializations and calculations, allowing you to see the relative performance of performing these algorithms sequentially and parallelized.

NOTE: Due to the number of classes that use ieng9 and the CPU-intensive nature of this assignment, we are developing and running the multi-threaded programming assignment (**PAMT1**) on the **workstations in the labs** or **ieng6** only! See the section on Getting the Code for more details on how to set up the pamt1 project on the workstations.

You will also be given a chance to practice your debugging skills. We've written a program that will read in strings from the command line, reverse the string, and determine whether each string is a palindrome (the same string forward and backward). We think it's close to working, but we didn't have time to debug it. It's your job to help us track down the bugs and fix it up so we can get it to work reversing strings. **The debugging exercise will be completed on ieng9.**

Important Points:

- 1) **PAMT1** will be developed and executed and turned in with your cs30x account on the **workstations in the lab** (preferred) or **ieng6.ucsd.edu** remotely.
- 2) **Debug1** will be debugged and executed and turned in with your cs30x account on **ieng9.ucsd.edu**.

Grading

PAMT1

- **README: 10 points** - See PAMT1 README File section
- **Compiling: 10 points** - Using our Makefile; no warnings. If what you turn in does not compile with the given Makefile, you will receive 0 points for this assignment. **NO EXCEPTIONS!**
- **Correctness: 40 points**

Debug1

- **README: 30 points** - See DEBUG README File section
- **Correctness: 10 points**

NOTE: If what you turn in does not compile with given Makefile, you will receive 0 points for this assignment.

PAMT1 Overview

For this assignment we will explore using OpenMP, an API that supports multi-platform shared memory multi-threading with minimal effort. Serial code can be easily made parallel by the addition of OpenMP directives (OpenMP pragmas) that will instruct the compiler to generate parallel code. This program will use a 12-byte encoded pattern that needs to be decoded to initialize an array of a user-specified length. Once the array is initialized correctly, the program will perform two modified sum calculations and find the min and max values. For each step of the project, we will perform it two different ways: 1) using a serial implementation and 2) with a parallel version. We will then compare the runtime of each version.

You will write three algorithms: `initData`, `squaredSumMinMax`, and `sqrtSumMinMax`. For each of these algorithms, you will implement a serial version and a parallel version. Each of the parallel versions will be identical to the corresponding serial version, but with added OpenMP pragmas. The project's `main()` is already written, the only modification that is required is commenting out five return statements (marked with TODOs) as you implement each of the algorithms. Each of these functions have already been stubbed out and function header comments provided, but the body must be added for correct functionality. You will also need to provide your personal info in the headers.

References

=====

<https://en.wikipedia.org/wiki/OpenMP>

- In particular, under "OpenMP clause" - pretty good intro to default sharing.
- Under "Data Sharing attribute clauses" - the explanation for "shared" and "private" are pretty good. Farther down the "Reduction" header (below "Data copying") is helpful.

<http://openmp.org/mp-documents/OpenMP3.1-CCard.pdf>

- First page/middle column - Parallel Loop shows syntax for `#pragma omp parallel for [clause]` and Simple Parallel Loop Example.
- Second page under "Clauses" - reduction table of operators (in our case: +, min, and max).

<http://openmp.org/wp/>

Detailed Overview

The function prototypes for the C functions

```
int main( int argc, char* argv[] );
void initData( unsigned char a[], size_t arraySize,
              unsigned char data[], size_t dataSize );
void parallel_initData( unsigned char a[], size_t arraySize,
                       unsigned char data[], size_t dataSize );
struct result squaredSumMinMax( unsigned char a[], size_t arraySize );
struct result parallel_squaredSumMinMax( unsigned char a[], size_t arraySize );
struct result1 sqrtSumMinMax( unsigned char a[], size_t arraySize );
struct result1 parallel_sqrtSumMinMax( unsigned char a[], size_t arraySize );
```

The definitions for `struct result` and `struct result1` are in the `pamt1.h` header file.

Getting the Code for PAMT1

Create a pamt1 directory in your cs30x class account on a workstation in the basement or on ieng6 (NOT in your cs30x class account while ssh'ed into ieng9!).

```
[cs30xzzz@ieng6-203]:~:420$ mkdir pamt1
```

Copy all the files in ~/../public/pamt1StarterFiles to your pamt1 directory.

```
[cs30xzzz@ieng6-203]:~:421$ cp ~/../public/pamt1StarterFiles/* ~/pamt1/
```

The files in your pamt1 directory should be:

Makefile	pamt1.h	parallel_squaredSumMinMax.c
initData.c	parallel_initData.c	sqrSumMinMax.c
main.c	parallel_sqrSumMinMax.c	squaredSumMinMax.c

1. initData.c

```
void initData( unsigned char a[], size_t arraySize,
               unsigned char data[], size_t dataSize );
```

This function will fill the array "a" of arraySize elements using the decoding algorithm listed below (and in the comments in the initData.c file) and the values in the "data" array of dataSize elements.

Steps:

1. Implement the decoding algorithm.

a) Be sure to use the constants defined in the pamt1.h header file to avoid use of magic numbers (BITS_IN_A_NIBBLE, BIT_7, BITS_IN_A_BYTE, etc.).

Decoding Algorithm:

- 1) Define 4 unsigned char variables b1, b2, b3, b4;
- 2) For size_t i from zero thru all the elements in array a
 - 3a) Set a[i] to data[i modulus dataSize]
 - 3b) Set b4 to zero
 - 3c) For int j from zero up to BITS_IN_A_NIBBLE (exclusive)
 - 4a) Set b1 to a[i] ANDed with (BIT_7 right shifted by j bits)
 - 4b) Set b1 to b1 right shifted by (BITS_IN_A_BYTE - 1 - j bits)
 - 4c) Set b2 to a[i] ANDed with (BIT_0 left shifted by j bits)
 - 4d) Set b2 to b2 right shifted by j bits
 - 4e) Set b3 to b1 XORed with b2
 - 4f) Set b3 to b3 left shifted by j bits
 - 4g) Set b4 to b4 ORed with b3
- 3d) Set a[i] to (a[i] ANDed with UPPER_NIBBLE) ORed with b4

2. Run 'make' in your pamt1 directory.

3. Run the program using './pamt1' followed by an array size (you can try multiple sizes to compare run-times).

a) The program will check that the first 12 decoded bytes in the "a" array match the correct pattern.

4. Once the function works correctly comment out the "return 0;" line with the TODO labeled "After initData (serial)" in main.c.

2. parallel_initData.c

```
void parallel_initData( unsigned char a[], size_t arraySize,  
                      unsigned char data[], size_t dataSize );
```

This function works identically to the serial version, but uses OpenMP pragmas to execute in parallel.

Steps:

1. Copy your working algorithm/code from the serial version of `initData()` (`initData.c`) into `parallel_initData.c`.

2. Add an OpenMP pragma to parallelize the outer for loop (iterating over the `arraySize`).

a) Adding OpenMP pragma.

In OpenMP every pragma begins with `"#pragma omp"` followed by the construct(s) we want to use. The pragma should be placed on the line preceding the block that it should apply to. We will add it on the line before the outer for loop.

b) Adding "parallel" construct.

In this case we will add the "parallel" construct to indicate that the next block of C code should be executed in multiple threads. The code in the parallel "region" (the for loop block in our case) will be duplicated and executed by each of the threads spawned (the number of threads spawned can vary, but the default is the total number of threads available on the compute node -- we will just use the default number of threads).

c) Adding "for" construct.

Here we don't want the code in the region to be duplicated and executed on each thread (this would execute the same for loop multiple times), instead we want each thread to execute a chunk of the iterations of the loop. To indicate this we will add the "for" construct after the "parallel" construct. Now the iterations of the for loop will be divided among the available threads (in chunks, by default). The line above the outer for loop should now look like this:

```
#pragma omp parallel for
```

3. Try re-compiling (`make`) and run the program. What happens?

a) You may notice that you are not getting the correct results for the parallel `initData`. This is most likely because of the scope of the unsigned char variables `b1`, `b2`, `b3` and `b4`. As stated in the wikipedia article on OpenMP (<https://en.wikipedia.org/wiki/OpenMP>): "By default, all variables in the work sharing region are shared except the loop iteration counter." In our case, the outer for loop block is our work sharing region.

4. Set the data scope attribute of the local variables.

a) In OpenMP, if a variable is defined inside a parallel region, then a unique copy of the variable will be used for each thread; this variable would be considered "private" to each thread. However, if a variable is defined outside of the region, it will be shared among all of the threads; this variable would be considered "shared".

b) The local variables `b1`, `b2`, `b3`, and `b4` are defined outside/above the outer for loop we are trying to parallelize, so by default they are shared between all the threads. The parallel threads will continuously overwrite each other's shared copies of `b1`, `b2`, `b3`, and `b4`. Not good and not what we want.

c) The solution (keeping the local variables defined at the top of the function above the for loop) is to add a "data scope attribute clause" to specify the scope. In our case we will add the "private" clause (e.g. `"private(b1, b2, b3, b4)"`). This indicates we want separate private copies of these variables for each thread of execution vs. having them all shared among all the threads.

5. Your OpenMP pragma should now look like:

```
#pragma omp parallel for private(b1, b2, b3, b4)
```

Note: If you defined the inner loop counter "j" as a local variable above the outer for loop, you will need to add "j" to the list of variables in the private clause. If you defined it as `int j = 0` in the inner for loop construct, then by default it is private in each outer loop work sharing region and you are fine.

6. Try re-compiling (make) and run the program. You should get a correct output for the parallel version that matches the serial version now. Use different array sizes to see if you can get the `initData Speed-up` to exceed 1.0, indicating that the parallel version ran faster than the serial version. You should see similar speed-up results to the example output, but they will probably not be exactly the same due to many factors such as other users/processes contending for time on the cores on the system you are running on.

7. Once the function works correctly comment out the "return 0;" line with the TODO labeled "After initData (parallel)" in `main.c`.

3. squaredSumMinMax.c

```
struct result squaredSumMinMax( unsigned char a[], size_t size );
```

This function will sum the squares of each element (square an element in the array and then add that value to an accumulating sum) in "a" and find the min and max values in the array.

The function returns a struct result. This struct has three members, unsigned long long "sum", unsigned int "min", and unsigned int "max". See `pamt1.h`.

Steps:

1. Create three local variables: sum, min and max with the same type as their correspondingly named struct result members. Use these during your calculations and set them into a struct result right before returning.

2. Initialize the local variables using the first element of the array (initialize sum to `a[0] * a[0]`, min and max to `a[0]`).

3. Iterate over all of the remaining elements in the array (starting with the second element).

a) Add the squared value of each element into the sum. Make sure to manually compute the square (instead of using the "pow" function, which returns a double).

b) Check each element to see if it is a new min or max and update the min and max accordingly.

4. Set the elements of a struct result using the local variables from your calculations and return the struct.

5. Try re-compiling (make) and run the program. Check your output against the provided example outputs for correctness.

6. Once the function works correctly comment out the "return 0;" line with the TODO labeled "After squaredSumMinMax (serial)" in `main.c`.

4. parallel_squaredSumMinMax.c

```
struct result parallel_squaredSumMinMax( unsigned char a[], size_t size );
```

This function works identically to the serial version, but uses OpenMP pragmas to execute in parallel.

Steps:

1. Copy your working algorithm/code from the serial version of squaredSumMinMax() in squaredSumMinMax.c into parallel_squaredSumMinMax.c.

2. Add an OpenMP pragma to parallelize the for loop.

a) Adding "parallel for" constructs.

Once again we will use the "parallel" and "for" constructs to parallelize the loop. Recompile the program and run. Compare the values of Squared Sum, Min, and Max between the sequential version and the parallel version of squaredSumMinMax. Hey, the Squared Sum values are different! Remember - by default, these local variables are being shared (except for the loop iteration counter variable).

b) Adding the "reduction" clause.

Now how will we define the data scope attribute of the sum, min and max variables? We do not want each thread overwriting the same shared variables; resulting in incorrect output. We also do not want to use a critical section (or atomics) to have all of the threads correctly share the variables because this will dramatically slow the program. We also do not want these variables to be strictly private to each thread (like we did for the local variables in parallel_initData.c) because in the end we want to combine each thread's local sum, min and max into a single overall sum, min, and max (this is called a reduction).

What we want is for each thread to have a private copy of each local variable during execution and then combine all of the threads' individual results in a shared value once they complete. To do this we will use the "reduction" clause. This clause will create a private copy of a variable for each thread and then combine the values in a shared variable using a specified operator once the threads finish (note that the private variables are initialized to a specific value based on the operator). Make sure that you are using local variables in the reduction clause, it will not work with struct members (e.g. "reduction(+:result.sum)" will not work).

1. For the "sum" we want to use "reduction(+:sum)".
2. For the "min" we want to use "reduction(min:min)".
3. For the "max" we want to use "reduction(max:max)".

3. Your OpenMP pragma should now look like:

```
#pragma omp parallel for reduction(+:sum) reduction(min:min) reduction(max:max)
```

or to put them on separate lines, use the backslash to escape the newline:

```
#pragma omp parallel for reduction(+:sum) \  
                        reduction(min:min) \  
                        reduction(max:max)
```

4. Try re-compiling and run the program. Make sure the parallel version's sum, min, and max values match those of the sequential version. Use different array sizes to see if you can get the squaredSumMinMax Speed-up to exceed 1.0, indicating that the parallel version ran faster. You should see similar speed-up results to the example output.

5. Once the function works correctly comment out the "return 0;" line with the TODO labeled "After squaredSumMinMax (parallel)" in main.c.

5. sqrtSumMinMax.c

```
struct result1 sqrtSumMinMax( unsigned char a[], size_t size );
```

This function will sum the square root of each element in "a" (call sqrt() on each element in the array and then add that value to an accumulating sum) and find the min and max values.

We will use the math library's "sqrt()" function for our calculations. See "man sqrt" for more information. Since this function returns a double, a separate "struct result1" with a double "sum" member is used as the return type.

Steps:

1. This function behaves almost exactly like squaredSumMinMax(). The only changes required are changing the data type of the "sum" variable and using sqrt() instead of squaring the array elements.
2. Try re-compiling and run the program. Check your output against the provided example outputs for correctness.
3. Once the function works correctly comment out the "return 0;" line with the TODO labeled "After sqrtSumMinMax (serial)" in main.c.

5. parallel_sqrtSumMinMax.c

```
struct result1 parallel_sqrtSumMinMax( unsigned char a[], size_t size );
```

This function works identically to the serial version, but uses OpenMP pragmas to execute in parallel.

Steps:

1. Using what you have learned about OpenMP, try adding the required pragma to parallelize this function.
2. Try re-compiling and run the program. Note that we are summing up double precision floating point values. Addition of floating point values is not associate, so we most likely will get different rounding errors between the sequential version and the parallel version - especially for large array sizes. But they should match that of the sample pamt1test. Use different array sizes to see if you can get the speed-up to exceed 1.0, indicating that the parallel version ran faster. You should see similar speed-up results to the example output.

There is a sample executable called pamt1test in ~/.public you can run. For example:

```
~/../public/pamt1test 40000000
```

Example Output

=====

```
[cs30xyz@ieng6]:pamt1:1$ ./pamt1 40000
```

```
Using array size = 40000
```

```
Initializing array with values using sequential initData()
```

```
[be patient with large values]
```

```
Checking that sequential initData produced the values as expected
```

```
Checking the first 12 bytes only
```

```
Should print out the string: "CSE30 Rocks!"
```

```
CSE30 Rocks!
```

```
Sequential initData time = 0.001554
```

```
Initializing array with same values using parallel initData()
[be less patient with large values]
Num of threads = 8
```

```
Checking that parallel initData produced the values as expected
Checking the first 12 bytes only
Should print out the string: "CSE30 Rocks!"
CSE30 Rocks!
```

```
Parallel initData time = 0.004342
```

```
*** initData Speed-up: 0.357939 ***
```

```
Sequential squared sum, min, max [be patient]
Sequential squaredSumMinMax time = 0.000191
```

```
Squared Sum is: 255583181
Min value is: 32
Max value is: 115
Completed in 0.000191 sec
```

```
Parallel squared sum, min, max [don't need to be as patient]
Parallel squaredSumMinMax time = 0.003774
```

```
Squared Sum is: 255583181
Min value is: 32
Max value is: 115
Completed in 0.003774 sec
```

```
*** squaredSumMinMax Speed-up: 0.050535 ***
```

```
Sequential sqrt sum, min, max [be patient]
Sequential sqrtSumMinMax time = 0.000590
```

```
Sqrt Sum is: 338939.708947
Min value is: 32
Max value is: 115
Completed in 0.000590 sec
```

```
Parallel sqrt sum, min, max [don't need to be as patient]
Parallel sqrtSumMinMax time = 0.002722
```

```
Sqrt Sum is: 338939.708947
Min value is: 32
Max value is: 115
Completed in 0.002722 sec
```

```
*** sqrtSumMinMax Speed-up: 0.216885 ***
```

```
-----
[cs30xyz@ieng6]:pamt1:2$ ./pamt1 40000000
Using array size = 40000000
Initializing array with values using sequential initData()
[be patient with large values]
```


Checking that sequential initData produced the values as expected
Checking the first 12 bytes only
Should print out the string: "CSE30 Rocks!"
CSE30 Rocks!

Sequential initData time = 1.508953

Initializing array with same values using parallel initData()
[be less patient with large values]
Num of threads = 8

Checking that parallel initData produced the values as expected
Checking the first 12 bytes only
Should print out the string: "CSE30 Rocks!"
CSE30 Rocks!

Parallel initData time = 0.295013

*** initData Speed-up: 5.114865 ***

Sequential squared sum, min, max [be patient]
Sequential squaredSumMinMax time = 0.173058

Squared Sum is: 255589993181
Min value is: 32
Max value is: 115
Completed in 0.173058 sec

Parallel squared sum, min, max [don't need to be as patient]
Parallel squaredSumMinMax time = 0.041531

Squared Sum is: 255589993181
Min value is: 32
Max value is: 115
Completed in 0.041531 sec

*** squaredSumMinMax Speed-up: 4.166964 ***

Sequential sqrt sum, min, max [be patient]
Sequential sqrtSumMinMax time = 0.586355

Sqrt Sum is: 338940858.017723
Min value is: 32
Max value is: 115
Completed in 0.586355 sec

Parallel sqrt sum, min, max [don't need to be as patient]
Parallel sqrtSumMinMax time = 0.117479

Sqrt Sum is: 338940858.042362
Min value is: 32
Max value is: 115
Completed in 0.117479 sec

*** sqrtSumMinMax Speed-up: 4.991132 ***

[cs30xyz@ieng6]:pamt1:3\$ **./pamt1 400000000**

Using array size = 400000000

Initializing array with values using sequential initData()
[be patient with large values]

Checking that sequential initData produced the values as expected
Checking the first 12 bytes only
Should print out the string: "CSE30 Rocks!"
CSE30 Rocks!

Sequential initData time = 15.239569

Initializing array with same values using parallel initData()
[be less patient with large values]
Num of threads = 8

Checking that parallel initData produced the values as expected
Checking the first 12 bytes only
Should print out the string: "CSE30 Rocks!"
CSE30 Rocks!

Parallel initData time = 2.649712

*** initData Speed-up: 5.751406 ***

Sequential squared sum, min, max [be patient]
Sequential squaredSumMinMax time = 1.733613

Squared Sum is: 2555899993181
Min value is: 32
Max value is: 115
Completed in 1.733613 sec

Parallel squared sum, min, max [don't need to be as patient]
Parallel squaredSumMinMax time = 0.365531

Squared Sum is: 2555899993181
Min value is: 32
Max value is: 115
Completed in 0.365531 sec

*** squaredSumMinMax Speed-up: 4.742720 ***

Sequential sqrt sum, min, max [be patient]
Sequential sqrtSumMinMax time = 5.892237

Sqrt Sum is: 3389408590.849287
Min value is: 32
Max value is: 115
Completed in 5.892237 sec

Parallel sqrt sum, min, max [don't need to be as patient]
Parallel sqrtSumMinMax time = 1.100185

```
Sqrt Sum is: 3389408590.272765
Min value is: 32
Max value is: 115
Completed in 1.100185 sec
```

```
*** sqrtSumMinMax Speed-up: 5.355679 ***
```

PAMT 1 README File

Along with your source code, you will be turning in a README (use all caps and no file extension for example, the file is README and not README.txt) file with every assignment. Use vi/vim to edit this file!

Your README file for this and all assignments should contain:

- Header with your name, cs30x login
- High level description of what your program does
- How to compile it (usually just typing "make")
- How to run it (give an example)
- An example of normal output and where that normal output goes (stdout or a file or ???)
- An example of abnormal/error output and where that error output goes (stderr usually)
- How you tested your program
- Anything else that you would want/need to communicate with someone who has not read the writeup
- Answers to questions (if there are any)

Questions to Answer for the README

The command `lscpu` displays information about the CPU(s) on that system. Running `lscpu` on one of the workstations in B240 shows there is a single socket (chip) with 4 cores per socket and 2 threads per core for a total of 8 logical CPUs. The 2 threads per core indicates the cores are hyper-threaded - sharing 2 threads of execution on a single physical CPU. Running `lscpu` on `ieng6.ucsd.edu` (say, `ieng6-201.ucsd.edu` or any of the other `ieng6-20[1-4]` front-end servers) shows they are configured with 8 sockets (chips) with 1 core per socket and 1 thread per core for a total of 8 logical CPUs. These are not hyper-threaded cores, so each core does not share its CPU resources with multiple threads.

Run `pamt1` with a large `array_size` (say, 400000000) on both `ieng6` and one of the workstations in B240.

1. Which system (`ieng6` or B240 workstation) shows the larger speed-up due to parallelization of the 3 algorithms?
2. Which system (`ieng6` or B240 workstation) runs faster (the different sequential and parallel completed times are smaller)?
3. Why do you think this is the case?
4. About what `array_size` values will result in the `initData Speed-up` to exceed 1.0? What about exceeding 2.0?
5. About what `array_size` values will result in the `squaredSumMinMax Speed-up` to exceed 1.0? What about exceeding 2.0?

6. About what array_size values will result in the sqrtSumMinMax Speed-up to exceed 1.0? What about exceeding 2.0?

Add an OpenMP parallel for pragma to also parallelize the inner for loop in parallel_initData.c.

7. How does this change the parallel speed-up? Does it make it faster or slower than not parallelizing the inner for loop? You can always comment out the omp parallel for pragma above the inner for loop to recompile and test with and without the inner for loop being parallelized.

8. Why do you think it is faster or slower?

Debugging Exercise Overview

Do this debugging exercise on ieng9.ucsd.edu

The purpose of this program is to read in all strings from the command line and find the reverse of that string. The main() function is written in C, and it will find the reverse of the string with help from several assembly functions.

We provide all of the code to you, but the code doesn't quite compile or work as it should. It is up to you to track down the bugs and fix them. There are a total of 8 bugs in the source code. You are required to record ALL of the bugs we placed and the solution for each bug. See the section on Debug README File for more details.

C routines

```
int main( int argc, char* argv[] );
```

Assembly routines

```
int reverse( char* str );  
int findEnd( char* str, char** endPtr );  
int swapChars( char* c1, char* c2 );
```

Getting the Code for Debugging Exercise

For debug1 (the debugging exercise), you will develop on **ieng9** as you do for most programming assignments. However, we will provide you with the buggy code. Simply go to your home directory and copy the whole folder to your home directory:

```
$ cd ~  
$ cp -r ~/../public/debug1 .
```

This will provide you with all of the buggy source code. All you have to do is fix the code and detail in a README file exactly what fixes you had to make to get this code to work properly - line numbers, what the line looked before and after your fix, etc. Be sure to include a short explanation of how you debugged each problem.

Debugging Exercise Example Output

The program takes one or more string arguments from the command line:

```
$ ./reverseString str1 [str2 str3 ...]
```

Each string will be printed to the screen, then reversed and printed to the screen again. If the string is a palindrome, the program will print a message saying so. At the end, the program prints the total number of palindromes found.

Below are a few examples (bold indicates user input):

```
[cs30xyz@ieng9]:pamt1$ ./reverseString potatoes
```

```
Before: potatoes
```

```
After: seotatop
```

```
You found 0 palindrome(s)
```

As you can see, the string entered on the command line is printed out, then reversed and printed out again. Let's see what happens if it's a palindrome...

```
[cs30xyz@ieng9]:pamt1$ ./reverseString amanaplanacanalpanama
```

```
Before: amanaplanacanalpanama
```

```
PALINDROME!
```

```
After: amanaplanacanalpanama
```

```
You found 1 palindrome(s)
```

This string is the same forward and backward, so we let the user know (with a triumphant PALINDROME!). We also print the number of palindromes found. (Note: we're not quite fancy enough to deal with spaces and punctuation, so the well-known palindrome "a man, a plan, a canal: Panama" won't work).

Now let's try entering several strings:

```
[cs30xyz@ieng9]:pamt1$ ./reverseString abba was a band with some serious wow  
factor
```

```
Before: abba
```

```
PALINDROME!
```

```
After: abba
```

```
Before: was
```

```
After: saw
```

```
Before: a
```

```
PALINDROME!
```

```
After: a
```

```
Before: band
```

```
After: dnab
```

Before: with
After: htiw

Before: some
After: emos

Before: serious
After: suoires

Before: wow
PALINDROME!
After: wow

Before: factor
After: rotcaf

You found 3 palindrome(s)

Aside from declaring my love for ABBA, this example shows what happens when several strings are entered on the command line, including some palindromes.

We can also enclose strings in quotes if we want to include spaces:

```
[cs30xyz@ieng9]:pamt1$ ./reverseString "I've always wanted to know how to spell  
my name in reverse"
```

Before: I've always wanted to know how to spell my name in reverse
After: esrever ni eman ym lleps ot woh wonk ot detnaw syawla ev'I

You found 0 palindrome(s)

```
[cs30xyz@ieng9]:pamt1$ ./reverseString "I was a saw I" "semolina is no meal" "---  
uuu-^U^-uuu---" "four score and seven years ago"
```

Before: I was a saw I
PALINDROME!
After: I was a saw I

Before: semolina is no meal
After: laem on si anilomes

Before: ---uuu-^U^-uuu---
PALINDROME!
After: ---uuu-^U^-uuu---

Before: four score and seven years ago
After: oga sraey neves dna erocs ruof

You found 2 palindrome(s)

Debugging C Modules

1. main.c

```
int main( int argc, char *argv[] );
```

The only C module of this program. Loops through all command line arguments in argv[]. It first prints the original string, then calls reverse() on it and prints the reversed string. If the string was a palindrome, a counter is incremented. When all strings have been read and reversed, a message is printed showing the total number of palindromes.

Return Value: Zero on success, nonzero on failure

Debugging Assembly Modules

1. reverse.s

```
int reverse( char* str );
```

The primary purpose of this function is to reverse the character array pointed to by str. It does this by finding the length of the string and a pointer to the last character of the string (using findEnd()) and then looping through all characters in the string, simultaneously incrementing the pointer at the front and decrementing the pointer at the back and swapping the characters. If the characters were the same (as returned by swapChars()), this function will keep track of that.

Return Value:

If all characters that were swapped were the same, this function will print a message ("PALINDROME!") and will return 1. Otherwise, it will not print any message and will return 0.

2. findEnd.s

```
int findEnd( char* str , char** endPtr );
```

This function has two purposes: to find the length of the string str and to set endPtr to point to the last character of the string. It does this simply by iterating through the string and checking whether the character is the null character, keeping a count of how many characters were seen. Once it finds the end of the string, it stores the pointer to the last character in endPtr.

Return Value:

Returns the length of the str.

3. swapChars.s

```
int swapChars( char* c1 , char* c2 );
```

Swaps the values of the two characters pointed to by c1 and c2. Determines if the characters were the same and, if so, if they were in fact the same character in the string (i.e. the addresses were the same).

Return Value:

If the characters were different, returns 0. If they were the same but the addresses were also the same, returns 1. If they were the same and the addresses were different, returns 2.

Debugging 1 README File

For the debugging assignment only, you do not have to include the usual high level description, how tested, etc. You will, however, have to list each of the compilation errors you encountered and the fix you made to correct the error (include the compilation error, the file name, the line number, and the new code).

You will also have to solve several runtime errors. Some of them will be obvious (for example, Bus Error), but some will involve a little more testing and debugging. Again, for each problem, describe the error and describe your solution for fixing it, including the file name, line number, and code for your fix.

As a guideline, there should be 2 compilation errors and 6 runtime/functionality problems. Make sure you locate all of them!! (Note: When we say there are 2 compilation errors, we mean that there are two fixes you'll have to make, not that there are two errors that are printed to the screen).

Turnin Instructions

Complete Turnin - due Friday night, May 13 @ 11:59 pm

Once you have checked your output, compiled, executed your code, and finished your README file (see below), you are ready to turn it in. Before you turn in your assignment, you should do make clean in order to remove all the object files, lint files, core dumps, and executables.

How to Turn in an Assignment

First, you need to have all the relevant files in a subdirectory of your home directory.

PAMT1

cse30turnin pamt1

In your cs30x account on the lab workstations of **ieng6**. You will not be able to turn-in pamt1 on ieng9.

Files required for PAMT1 Final Turn-in:

initData.c	parallel_initData.c	squaredSumMinMax.c
main.c	sqrtSumMinMax.c	parallel_squaredSumMinMax.c
pamt1.h	parallel_sqrtSumMinMax.c	

DEBUG1

cse30turnin debug1

In your cs30x account on **ieng9**. You will not be able to turn-in debug1 on ieng6.

Files required for Debug1 Final Turn-in:

debug.h	main.c	swapChars.s
findEnd.s	reverse.s	

Style Requirements

You will be graded for the style of programming on all the assignments. A few suggestions/requirements for style are given below. Read carefully, and if any of them need clarification do not hesitate to ask.

- Use reasonable comments to make your code clear and readable.
- Use file headers and function header blocks to describe the purpose of your programs and functions. Sample file/function headers are provided with PA0.

- Explicitly comment all the various registers that you use in your assembly code.
- In the assembly routines, you will have to give high level comments for the synthetic instructions, specifying what the instruction does.
- You should test your program to take care of invalid inputs like nonintegers, strings, no inputs, etc. This is very important. Points will be taken off if your code doesn't handle exceptional cases of inputs.
- Use reasonable variable names.
- Error output goes to stderr. Normal output goes to stdout.
- Use `#defines` and assembly constants to make your code as general as possible.
- Use a local header file to hold common `#defines`, function prototypes, type definitions, etc., but not variable definitions.
- Judicious use of blank spaces around logical chunks of code makes your code much easier to read and debug.
- Keep all lines less than 80 characters, split long lines if necessary.
- Use 2-4 spaces for each level of indenting in your C source code (do not use tab). Be consistent. Make sure all levels of indenting line up with the other lines at that level of indenting.
- Do use tabs in your Assembly source code.
- Always recompile and execute your program right before turning it in just in case you commented out some code by mistake.
- Before running turnin please do a make clean in your project directory.
- Do `#include` only the header files that you need and nothing more.
- Always macro guard your header files (`#ifndef ... #endif`).
- Never have hardcoded magic numbers. This means we shouldn't see magic constants sitting in your code. Use a `#define` if you must instead.